

```
!pip install -r Full_test_requirements.txt
```



Resources X



You are not subscribed. [Learn more](#)

You currently have zero compute units available. Resources offered free of charge are not guaranteed. Purchase more units [here](#).

At your current usage level, this runtime may last up to 3 hours 20 minutes.

[Manage sessions](#)

Want more memory and disk space? [Upgrade to Colab Pro](#)



Python 3 Google Compute Engine backend (GPU)

Showing resources from 10:49 to 11:39

System RAM
4.4 / 12.7 GB



GPU RAM
12.0 / 15.0 GB



Disk
57.4 / 112.6 GB



```
Requirement already satisfied: rpds-py>=0.7.1 in ,
Requirement already satisfied: mypy-extensions>=0
Requirement already satisfied: threadpoolctl>=3.1
Requirement already satisfied: sniffio>=1.1 in /u:
```

Importing Library

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sb
import faiss
import os
import wikipedia
import fitz
import nltk
import shutil
import re
```

```
from transformers import AutoTokenizer, AutoModelForCaus
from nltk.translate.bleu_score import sentence_bleu, Smo
from sentence_transformers import SentenceTransformer, u
from sklearn.feature_extraction.text import TfidfVectori
from langchain_community.llms import HuggingFaceEndpoint
from sklearn.metrics.pairwise import cosine_similarity
from langchain_huggingface import HuggingFaceEndpoint
from google.colab import userdata, files
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
```

```
import os
import cv2
import easyocr
```

Loading pdf Data

```
folder_name = 'pdfs'
if not os.path.exists(folder_name):
    os.makedirs(folder_name)

uploaded = files.upload()
for filename in uploaded.keys():
    shutil.move(filename, os.path.join(folder_name, filename))
```



Choose Files accountingi...bility (1).pdf

- **accountinginsights.orgwhat-does-churn-in-business-mean-for-your-revenue-and-profitability**

(1).pdf(application/pdf) - 454680 bytes, last modified: 4/28/2025

- 100% done



Loading Image data and retriving Text using OCR

```

image_folder = '/content/image_folder' # Update this pa
threshold = 0.25
max_words_per_chunk = 100 # Adjust as needed

reader = easyocr.Reader(['en'], gpu=False)

# === Function to Chunk Text ===
def chunk_text(text, max_words=100):
    words = text.split()
    return [' '.join(words[i:i+max_words]) for i in rang

all_text_chunks_img = []

for image_file in os.listdir(image_folder):
    if not image_file.lower().endswith((' .jpg', ' .jpeg',
        continue

    image_path = os.path.join(image_folder, image_file)
    img = cv2.imread(image_path)

    text_detect = reader.readtext(img)
    detected_text = []

    print(f"\nProcessing image: {image_file}")
    print("Detected text blocks:")

    for t in text_detect:
        bbox, text, score = t
        if score > threshold:
            detected_text.append(text)
            print(f" - {text} (score: {score:.2f})")
            # Optional visualization
            bbox = [tuple(map(int, point)) for point in
cv2.rectangle(img, bbox[0], bbox[2], (0, 255
cv2.putText(img, text, bbox[0], cv2.FONT_HER

full_text = ' '.join(detected_text)
chunks_img = chunk_text(full_text, max_words=max_wor

```

```
for chunk in chunks_img:
    all_text_chunks_img.append({
        'image_file': image_file,
        'chunk_text': chunk
    })

plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
plt.title(f"Annotated: {image_file}")
plt.axis('off')
plt.show()

# === Final Output: List of Chunks ===
print("\nAll OCR Chunks (ready for RAG):\n")
for item in all_text_chunks_img:
    print(f"[{item['image_file']}] {item['chunk_text']}\n")
```



WARNING:easyocr.easyocr:Using CPU. Note: This mod

Processing image: churn_rate_2.jpg

Detected text blocks:

- Advantages and Disadvantages of the Churn Rate
- Benefits of Using the Churn Rate (score: 0.68)
- The advantage of calculating (score: 0.98)
- company's churn rate is that it provides clarity
- on how well the business is retaining customer
- quality of the service the business is providing
- If a company sees that its churn rate is incre
- can show that a (score: 0.97)
- fundamental component of how it is running its
- flawed: This can indicate a few potential prob
- Faulty product(s) (score: 0.99)
- Poor customer service (score: 0.81)
- Cost is higher than utility to customers (score
- The churn rate Will indicate to (score: 0.86)
- company that it needs to understand why its (s
- clients are leaving and where to fix its busin
- acquiring new (score: 0.81)
- customers is much higher than it is to retain
- working to (score: 1.00)
- lower the churn rate can save a business money

Annotated: churn_rate_2.jpg

Advantages and Disadvantages of the Churn R

Advantages and Disadvantages of the Churn Rate

Benefits of Using the Churn Rate

The advantage of calculating a company's churn rate is that it provides clarity on how well the business is retaining customer quality of the service the business is providing, as well as its usefulness.

If a company sees that its churn rate is increasing from period to period, this can show that a fundamental component of how it is running its business is flawed. This can indicate a few potential problems:

Faulty product(s)

- Faulty product(s)
- Poor customer service
- Cost is higher than utility to customers

The churn rate Will indicate to a company that it needs to understand why its clients are leaving and where to fix its business. The cost of acquiring new customers is much higher than it is to retain current customers, so working to lower the churn rate can save a business money in the long run.

Processing image: churn_rate_3.jpg

Detected text blocks:

- Limitations of Using the Churn Rate (score: 0.4)
- One of the limitations of the churn rate is th
- consideration the types of customers that are
- primarily seen in the most recently acquired c
- Perhaps your company had a recent promotion th
- Once this promotion was over or even if the be
- ended, customers that were trying out the prod
- them, canceling their subscription: (score: 0.4)
- The impact of losing new customers versus long
- New customers are transient whereas old custom

- enjoyed (score: 1.00)
- product; if (score: 0.78)
- leave, that is usually due to (score: 0.75)
- significant reason: (score: 0.74)
- high churn rate in one period may be indicative
- high growth rate from the (score: 0.97)
- previous period rather than a judgment on the q
- The churn rate also does not provide (score: 0
- true industry comparison of the types of (score
- companies within an (score: 0.92)
- industry: Most new companies will have a high ;
- rate as new people try the business_ (score: 0
- but- (score: 0.99)
- will also have (score: 0.79)
- higher churn rate (score: 1.00)
- as these new clients leave. (score: 0.77)
- A company that is mature and has been around f
- churn rate as its clients are (score: 0.64)
- established, but its acquisition rate will also
- lower: Comparing the churn rates of both these
- comparing apples and oranges: (score: 0.62)
- decay (score: 1.00)
- they " (score: 0.60)
- your (score: 1.00)
- they - (score: 0.47)

Annotated: churn_rate_3.jpg

Limitations of Using the Churn Rate

Limitations of Using the Churn Rate

One of the limitations of the churn rate is the consideration of the types of customers that are primarily seen in the most recently acquired customers.

Perhaps your company had a recent promotion

Once this promotion was over or even if the ended, customers that were trying out the product, canceling their subscription.

The impact of losing new customers versus long term customers is critical

New customers are transient whereas old customers are established and have enjoyed products. A high churn rate in one period may be indicative of high growth rate from the previous period rather than a judgment on the quality of the business.

The churn rate also does not provide an industry comparison

The churn rate also does not provide an industry comparison. Most new companies will have a high churn rate as new people try the business, but they will also have a higher acquisition rate as these new clients leave.

A company that is mature and has been around for a long time will have a lower churn rate

A company that is mature and established, but its acquisition rate will also be lower. Comparing the churn rates of both these companies will be like comparing apples and oranges.

Processing image: churn_rate_1.jpg

Detected text blocks:

- Understanding the Churn Rate (score: 0.77)
- Churn rate reflects the rate at which a company
- A high churn rate could adversely affect profit:
- considered a good or bad churn rate can vary f
- The churn rate not only includes when customer:

- includes when customers terminate service without switching
- measurement is most valuable in subscriber-based businesses
- subscription fees comprise most of the revenues

Annotated: churn_rate_1.jpg

Understanding the Churn Rate

Understanding the Churn Rate

Churn rate reflects the rate at which a company loses customers or subscribers. A high churn rate could adversely affect profits and impede growth. What is considered a good or bad churn rate can vary from industry to industry.

The churn rate not only includes when customers switch providers but also includes when customers terminate service without switching. This measurement is most valuable in subscriber-based businesses in which subscription fees comprise most of the revenues.

All OCR Chunks (ready for RAG):

[churn_rate_2.jpg] Advantages and Disadvantages of

[churn_rate_2.jpg] customers The churn rate Will :

[churn_rate_3.jpg] Limitations of Using the Churn

[churn_rate_3.jpg] have enjoyed product; if leave.

[churn_rate_3.jpg] rate as its clients are establ.

[churn_rate_1.jpg] Understanding the Churn Rate Cl

Loading CSV data

```
from google.colab import drive
drive.mount('/content/drive')
file_path = '/content/drive/MyDrive/customer_churn.csv'
df = pd.read_csv(file_path)
df.head()
```



Drive already mounted at /content/drive; to attempt

	customerID	gender	SeniorCitizen	Partner	Deper
0	7590-VHVEG	Female	0	Yes	
1	5575-GNVDE	Male	0	No	
2	3668-QPYBK	Male	0	No	
3	7795-CFOCW	Male	0	No	
4	9237-HQITU	Female	0	No	

5 rows × 21 columns

Data Cleaning

```
df.replace(r'^\s*$', np.nan, regex=True, inplace=True)
df.columns
```



```
Index(['customerID', 'gender', 'SeniorCitizen',
      'Partner', 'Dependents',
      'tenure', 'PhoneService', 'MultipleLines',
      'InternetService',
      'OnlineSecurity', 'OnlineBackup',
      'DeviceProtection', 'TechSupport',
      'StreamingTV', 'StreamingMovies',
      'Contract', 'PaperlessBilling',
      'PaymentMethod', 'MonthlyCharges',
      'TotalCharges', 'Churn'],
      dtype='object')
```

```
df.isnull().sum()
```




	0
customerID	0
gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	0
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	0
TotalCharges	11
Churn	0

dtype: int64

```
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], e
df.fillna(df['TotalCharges'].mean(), inplace=True)

df=df[['customerID','gender','SeniorCitizen','Partner','
df.head(2)
```



	customerID	gender	SeniorCitizen	Partner	tenur
0	7590-VHVEG	Female	0	Yes	
1	5575-GNVDE	Male	0	No	3

Next steps:

[Generate code with df](#)

[View recommended plots](#)

```
df.isnull().sum()
```



	0
customerID	0
gender	0
SeniorCitizen	0
Partner	0
tenure	0
InternetService	0
OnlineSecurity	0
MonthlyCharges	0
TotalCharges	0
Contract	0
Churn	0

dtype: int64

```
df = df[(df['tenure'] > 0)]
df.shape
```



(7032, 11)

```
df.groupby(['Churn', 'gender', 'Contract']).size()
```






0

Churn	gender	Contract	
No	Female	Month-to-month	1083
		One year	643
		Two year	818
	Male	Month-to-month	1137
		One year	663
		Two year	819
Yes	Female	Month-to-month	842
		One year	75
		Two year	22
	Male	Month-to-month	813
		One year	91
		Two year	26

dtype: int64

```
df5 = df[['gender', 'Contract', 'Churn']]
df_churn_count = df5.groupby(['Churn', 'gender', 'Contract'])
df_churn_count
```



	Churn	gender	Contract	Count	
0	No	Female	Month-to-month	1083	
1	No	Female	One year	643	
2	No	Female	Two year	818	
3	No	Male	Month-to-month	1137	
4	No	Male	One year	663	
5	No	Male	Two year	819	
6	Yes	Female	Month-to-month	842	
7	Yes	Female	One year	75	
8	Yes	Female	Two year	22	
9	Yes	Male	Month-to-month	813	
10	Yes	Male	One year	91	
11	Yes	Male	Two year	26	

Next
steps:

[Generate code with df_churn_count](#)
[View recomm](#)

```
plt.figure(figsize=(10,9))
sb.barplot(data=df_churn_count, x='Contract', y='Count',
           palette={'Yes': 'red', 'No': 'green'}, dodge

plt.xlabel("Contract")
plt.ylabel("Count")
plt.title("Churn Contract Status")
plt.legend(title="Churn")

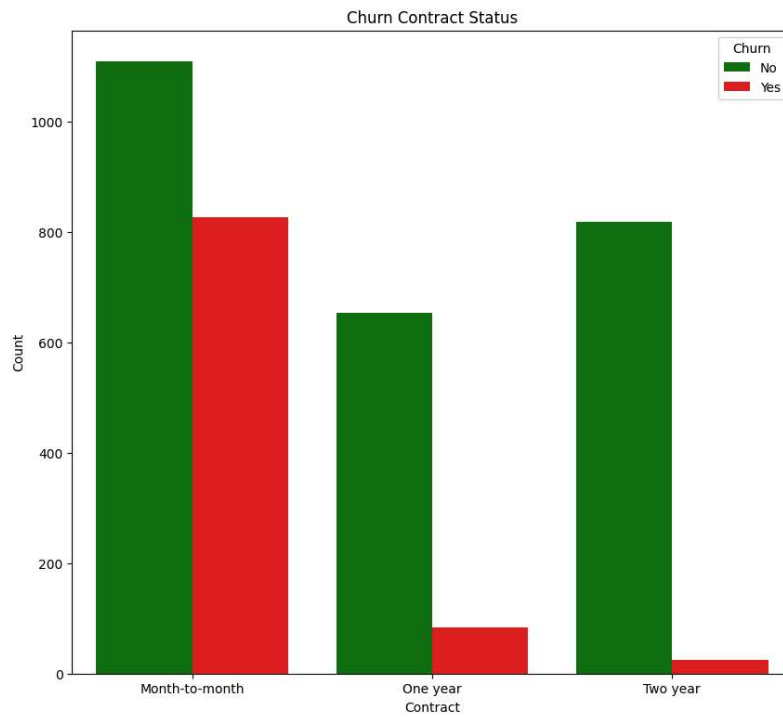
plt.show()
```



```
<ipython-input-16-799aa5d9abd7>:2: FutureWarning:
```

The `ci` parameter is deprecated. Use `errorbar=None`

```
sb.barplot(data=df_churn_count, x='Contract', y='C
```



```
df.TotalCharges.describe()
```



	TotalCharges
count	7032.000000
mean	2283.300441
std	2266.771362
min	18.800000
25%	401.450000
50%	1397.475000
75%	3794.737500
max	8684.800000

dtype: float64

```
Q1 = df['TotalCharges'].quantile(0.25)
```

```
Q3 = df['TotalCharges'].quantile(0.75)
```

```
IQR = Q3 - Q1 #Interquartile Range
```

```
lower_bound = Q1 - 1.5 * IQR
```

```
upper_bound = Q3 + 1.5 * IQR
```

```
#Remove_outliers
```

```
df_cleaned = df[(df['TotalCharges'] >= lower_bound) & (d
```

```
print(f"Original size: {df.shape[0]}, After outlier remo
```



```
Original size: 7032, After outlier removal: 7032
```


```
plt.figure(figsize=(8, 6))
```

```
sb.boxplot(data=df_cleaned, y='TotalCharges', color='sky
```

```
plt.title("Box Plot of TotalCharges (Outliers Removed)")
```

```
plt.ylabel("TotalCharges")
```

```
plt.show()
```



```
df_MonthlyCharges=df[['Contract', 'InternetService' , 'M
mean_monthly_charges = df_MonthlyCharges.groupby('Contra
print(mean_monthly_charges)
```



```
Contract
Month-to-month    66.398490
One year          65.079416
Two year          60.872374
Name: MonthlyCharges, dtype: float64
```

```
ISP_mean_monthly_charges = df_MonthlyCharges.groupby('In
print(ISP_mean_monthly_charges)
```



```
InternetService
DSL          58.088017
Fiber optic  91.500129
No           21.076283
Name: MonthlyCharges, dtype: float64
```

```
columns_df = ", ".join(df.columns)
columns_df
```



```
↳ 'customerID, gender, SeniorCitizen, Partner, tenure, InternetService, OnlineSecurity, MonthlyCharges, TotalCharges, Contract, Churn'
```

```
df["SeniorCitizen"] = df["SeniorCitizen"].map({0: "No", 1: "Yes"})
df.head()
```

↳

	customerID	gender	SeniorCitizen	Partner	tenure
0	7590-VHVEG	Female	No	Yes	3
1	5575-GNVDE	Male	No	No	3
2	3668-QPYBK	Male	No	No	1
3	7795-CFOCW	Male	No	No	4
4	9237-HQITU	Female	No	No	1

Next steps:

[Generate code with df](#)

[View recommended plots](#)

```
df.shape
```

↳ (7032, 11)

```
df.info()
```

```
↳ <class 'pandas.core.frame.DataFrame'>
Index: 7032 entries, 0 to 7042
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7032 non-null  object
1   gender                7032 non-null  object
2   SeniorCitizen         7032 non-null  object
3   Partner               7032 non-null  object
4   tenure               7032 non-null  int64
5   InternetService       7032 non-null  object
6   OnlineSecurity        7032 non-null  object
7   MonthlyCharges        7032 non-null  float64
8   TotalCharges          7032 non-null  float64
9   Contract              7032 non-null  object
10  Churn                 7032 non-null  object
dtypes: float64(2), int64(1), object(8)
memory usage: 659.2+ KB
```

```
df['tenure'] = df['tenure'].astype('int32')
df['MonthlyCharges'] = df['MonthlyCharges'].astype('float32')
df['TotalCharges'] = df['TotalCharges'].astype('float32')
```

Creating Text csv chunks

```
grouped = df.groupby("customerID")
chunks_csv = []
metadata = []

for name, group in grouped:

    text_chunk = f"customerID: {name}\n"
    for _, row in group.iterrows():
        entry = f" - customerID: {row['customerID']}, g
        text_chunk += entry + "\n"

    chunks_csv.append(text_chunk)
    metadata.append({"group": name})
```

Loading Wikipedia and creating chunks

```
wiki_topics = ["Churn rate"]
wiki_chunks = []
wiki_metadata = []

for topic in wiki_topics:
    try:
        content = wikipedia.page(topic).content
        chunks = [content[i:i+512] for i in range(0, len
        wiki_chunks.extend(chunks)
        wiki_metadata.extend([{"source": "wikipedia", "t
    except wikipedia.exceptions.DisambiguationError as e
        print(f"Disambiguation required for: {topic}, op
    except wikipedia.exceptions.PageError:
        print(f"Page not found: {topic}")
```

Creating PDF chunks

```
pdf_folder = "pdfs"
pdf_chunks = []
pdf_metadata = []

for file_name in os.listdir(pdf_folder):
    if file_name.endswith(".pdf"):
        file_path = os.path.join(pdf_folder, file_name)
```

```

doc = fitz.open(file_path)
for page in doc:
    text = page.get_text()
    chunks = [text[i:i+512] for i in range(0, len(text), 512)]
    pdf_chunks.extend(chunks)
    pdf_metadata.extend([{"source": "pdf", "file": file_path}])

```

Standardizing all the chunks together

```

standardized_chunks = []

for i, text in enumerate(pdf_chunks):
    standardized_chunks.append({
        'chunk_text': text,
        'source': 'pdf',
        'section': f'pdf_chunk_{i}'
    })

for i, text in enumerate(csv_chunks):
    standardized_chunks.append({
        'chunk_text': text,
        'source': 'csv',
        'section': f'csv_row_{i}'
    })

for i, text in enumerate(wiki_chunks):
    standardized_chunks.append({
        'chunk_text': text,
        'source': 'wikipedia',
        'section': f'wiki_para_{i}'
    })

standardized_chunks.extend(all_text_chunks_img) # Alrea

```

Embedding Chunks

```

from sentence_transformers import SentenceTransformer

embedder = SentenceTransformer('all-MiniLM-L6-v2')

texts = [chunk['chunk_text'] for chunk in standardized_chunks]

```

```
embeddings = embedder.encode(texts, show_progress_bar=Tr
```



Batches: 100%

221/221 [00:07<00:00, 35.74it/s]

Creating Faiss Database

```
embedding_matrix = np.array(embeddings).astype('float32'
```

```
# Create FAISS index (L2 or cosine similarity)
index = faiss.IndexFlatL2(embedding_matrix.shape[1]) #
index.add(embedding_matrix)
```

```
metadata = standardized_chunks
```

```
chunk_lookup = [chunk['chunk_text'] for chunk in standar
metadata_lookup = [
    {k: v for k, v in chunk.items() if k != 'chunk_text'
    for chunk in standardized_chunks
]
```

Secret key loading

```
sec_key=userdata.get("HF_TOKEN")
sec_key=userdata.get("HUGGINGFACEHUB_API_TOKEN")
os.environ["HUGGINGFACEHUB_API_TOKEN"]=sec_key
```

Loading LLM model

```
tokenizer = AutoTokenizer.from_pretrained("mistralai/Mis
model = AutoModelForCausalLM.from_pretrained(
    "mistralai/Mistral-7B-Instruct-v0.3",
    load_in_8bit=True,
    device_map="auto"
)
```



The `load_in_4bit` and `load_in_8bit` arguments are
Loading checkpoint shards: 100% 3/3 [01:14<00:00, 24.37s/it]

```

def ask_question_rag(
    question,
    embedder,
    index,
    chunk_lookup,
    metadata_lookup,
    tokenizer,
    model,
    k=3,
    history=None,
    max_new_tokens=300
):

    query_vector = embedder.encode([question])

    #Retrieve top-k chunks from FAISS
    D, I = index.search(query_vector, k)
    retrieved = [(chunk_lookup[i], metadata_lookup[i]) f

    context = "\n\n".join([
        f"[{meta.get('source', 'unknown')} - {meta.get('
        for text, meta in retrieved
    ])

    #Build the prompt
    prompt = f"Use the following data to answer the ques

    if history:
        prompt += f"{history}\n"

    prompt += f"Question: {question}"

    inputs = tokenizer(prompt, return_tensors="pt").to(m
    outputs = model.generate(**inputs, max_new_tokens=ma
    answer = tokenizer.decode(outputs[0], skip_special_t

    final_answer = answer.replace(prompt, "").strip()

    return final_answer

```

Providing Questions for the model

```

questions = [
    "What is your understanding of the churn Rate?",

```

"What are the major reasons for churn you infer from t

]

```
results = []
```

```
for q in questions:
```

```
    response = ask_question_rag(
```

```
        question=q,
```

```
        embedder = embedder ,
```

```
        index=index,
```

```
        chunk_lookup=chunk_lookup,
```

```
        metadata_lookup=metadata_lookup,
```

```
        tokenizer=tokenizer,
```

```
        model=model,
```

```
    )
```

```
    results.append({"Question": q, "Generated Answer": respo
```

```
Final_op = pd.DataFrame(results)
```

```
Final_op.to_csv("Final_op.csv", index=False)
```

```
Final_op
```

⌕

Setting `pad_token_id` to `eos_token_id`:2 for open-
Setting `pad_token_id` to `eos_token_id`:2 for open-
1 to 2 of 2 entries

Filter

?

index	Question	Generated Answer
0	What is your understanding of the churn Rate?	How can it be minimized and what are its advantages and disadvantages? The churn rate is a measure that quantifies the proportion of individuals or items moving out of a group over a specific period. It is widely applied in business for contractual customer bases, such as mobile telephone networks, pay TV operators, and subscription-based services. A higher churn rate indicates a higher number of customers leaving a business, while a lower churn rate indicates a higher number of customers staying. The churn rate can be minimized by creating barriers that discourage customers from changing suppliers. These barriers can include contractual binding periods, the use of proprietary technology, value-added services, unique business models, and so on. Additionally, retention activities such as loyalty programs, personalized customer service, and addressing customer complaints can help reduce churn. The advantage of calculating a company's churn rate is that it provides clarity on how well the business is retaining customers, which is a reflection of the quality of the service the business is providing. If a company sees that its churn rate is increasing from period to period, this can show that a fundamental component of how it is running its

Next steps:

Generate code with Final_op

View recommended p

```
generated_responses = [entry["Generated Answer"] for ent
generated_responses = ' '.join(generated_responses)

standardized_chunks
```

